

## **BCM963XX**

### **CPE Linux Ingress QoS**

#### **Application Note**

Broadcom Confidential

---

Broadcom, the pulse logo, Connecting everything, Avago Technologies, Avago, and the A logo are among the trademarks of Broadcom and/or its affiliates in the United States, certain other countries, and/or the EU.

Copyright © 2010–2018 Broadcom. All Rights Reserved.

The term “Broadcom” refers to Broadcom Inc. and/or its subsidiaries. For more information, please visit [www.broadcom.com](http://www.broadcom.com).

Broadcom reserves the right to make changes without further notice to any products or data herein to improve reliability, function, or design. Information furnished by Broadcom is believed to be accurate and reliable. However, Broadcom does not assume any liability arising out of the application or use of this information, nor the application or use of any product or circuit described herein, neither does it convey any license under its patent rights nor the rights of others.

Broadcom Confidential

---

# Table of Contents

<b>Chapter 1: Overview .....</b>	<b>5</b>
1.1 Purpose and Audience .....	5
1.2 Acronyms and Abbreviations .....	5
<b>Chapter 2: Enable and Disable Ingress QoS .....</b>	<b>6</b>
2.1 Enable Ingress QoS .....	6
2.2 Disable Ingress QoS .....	6
<b>Chapter 3: Ingress QoS .....</b>	<b>7</b>
3.1 CPE Behavior — No Congestion State .....	7
3.2 CPE Behavior — Congestion State .....	7
3.3 Ingress QoS Packet Priority .....	8
3.4 Force Trap .....	8
3.5 Drop .....	9
<b>Chapter 4: CLI Commands .....</b>	<b>10</b>
4.1 Status .....	11
4.2 Enable .....	12
4.3 Disable .....	12
4.4 Flush .....	12
4.5 Add Port .....	13
4.6 Remove Port .....	13
4.7 Get Port .....	13
4.8 Dump Key Mask and Key Table .....	14
4.9 Add Key Mask .....	16
4.10 Remove Key Mask .....	16
4.11 Dump Key Mask Table .....	17
4.12 Add Key .....	17
4.13 Remove Key .....	18
4.14 Get Key .....	18
4.15 Dump Key Table .....	19
4.16 Set Default Priority for IP Protocol Value .....	20
4.17 Remove Configured Default Priority of an IP Protocol Value .....	20
4.18 Example .....	20
4.19 Example .....	21
<b>Chapter 5: APIs .....</b>	<b>22</b>
5.1 Add L4 Destination Port .....	22
5.2 Remove L4 Destination Port .....	22
5.3 Get L4 Destination Port Priority .....	22
5.4 Initialize Keymask Parameter .....	23

5.5 Set Field and Value to the Keymask Parameter .....	23
5.6 Commit and Add the Keymask Parameter .....	23
5.7 Commit and Delete the Keymask Parameter .....	23
5.8 Initialize Key Parameter .....	24
5.9 Set Field and Value to the Key Parameter .....	24
5.10 Set Action and Value to the Key Parameter .....	24
5.11 Commit and Add the Key Parameter .....	24
5.12 Commit and Delete the Key Parameter .....	25
5.13 Commit and Get the Key Parameter .....	25
5.14 Flush All Dynamic Key Entries .....	25
5.15 Set Status .....	26
5.16 Configuration Sequence .....	27
<b>Chapter 6: Performance .....</b>	<b>28</b>
6.1 All High-Priority Flows .....	28
6.2 All Low-Priority Flows .....	28
6.3 Mix of High- and Low-Priority Flows .....	28
6.4 Early Trap and Drop .....	28
<b>Revision History .....</b>	<b>29</b>
963XX-AN403-R, December 18, 2018 .....	29
963XX-AN402-R, July 11, 2018 .....	29
Previous Release History .....	30

# Chapter 1: Overview

The Ingress QoS feature provides a central control unit for managing different Ingress QoS needs, such as:

- Providing a mechanism to allow high priority traffic to be prioritized over lower priority traffic when there is CPU and/or packet processing resource congestion.
- Interacting with a hardware network processor to place packets in high priority receive queue, trap the packets in early packet classification stage, or drop the packets without informing the host processor.

**NOTE:** Within the scope of this document, “Congestion” means “CPU and/or packet processing resource congestion.” Both terms are used interchangeably.

This document is organized into the following sections:

- “[Enable and Disable Ingress QoS](#)” explains how to enable or disable the Ingress QoS feature.
- “[Ingress QoS](#)” explains the Ingress QoS features in detail.
- “[CLI Commands](#)” describes the Ingress QoS CLI commands.
- “[APIs](#)” describes various Ingress QoS APIs.
- “[Performance](#)” describes the impact of the Ingress QoS feature on low, high, or a combination of low and high priority traffic.

## 1.1 Purpose and Audience

The Ingress QoS feature allows high priority traffic to be prioritized when there is CPU congestion or resource limitations. This document describes the Ingress QoS feature from a user perspective, and is aimed at engineers using the BCM963XX reference design boards or designing with BCM63XX chips.

## 1.2 Acronyms and Abbreviations

In most cases, acronyms and abbreviations are defined on first use. Acronyms and abbreviations in this document are shown in [Table 1](#).

**Table 1: Abbreviations**

<b>Abbreviation</b>	<b>Definition</b>
CPE	Customer Premises Equipment
IP	Internet Protocol version 4
IQ	Ingress QoS
L4	Layer 4 protocol (TCP/UDP)
QoS	Quality of Service
RXBD	RX buffer descriptor
TCP	Transmission Control Protocol
UDP	User Datagram Protocol

## Chapter 2: Enable and Disable Ingress QoS

By default, the Ingress QoS feature is enabled in all of the profiles.

**NOTE:** The Ingress QoS feature can be either statically built (\*) with the Linux kernel or built as module (M) or be compiled out.

### 2.1 Enable Ingress QoS

To enable the Ingress QoS feature:

1. Use the `make menuconfig` command at the Linux command prompt before build.  
`$ make menuconfig`
2. From **menuconfig**, select **Datapath**, then **Buffer Pool Manager and Ingress QoS**, and then **Ingress QoS**.
3. Use the space bar to select the Ingress QoS feature.

### 2.2 Disable Ingress QoS

To disable the Ingress QoS feature:

1. Use the `make menuconfig` command at the Linux command prompt before build.  
`$ make menuconfig`
2. From **menuconfig**, select **Datapath**, then **Buffer Pool Manager and Ingress QoS**, and then **Ingress QoS**.
3. Use the space bar to deselect the Ingress QoS feature.

## Chapter 3: Ingress QoS

The Ingress QoS feature protects high-priority traffic under CPU congestion or resource limitations by causing an early drop of low-priority traffic.

CPU and/or other packet processing resource congestion determination is platform dependent and kept outside the scope of this document.

Notes:

- Protection of high-priority packets means processing of all high-priority packets without drop subject to egress bandwidth limitations
- Packet priority (low or high) refers strictly to Ingress QoS packet priority and does not refer to any other priority like IEEE 802.1p, IPv4 ToS, etc. However, one can configure Ingress QoS based on any of these other priorities.
- Ingress QoS refers to QoS at the ingress/RX interface. This is not an end-to-end QoS (from RX to TX interface) feature. For an end-to-end QoS, Ingress QoS should be used along with the egress QoS feature to prioritize the flows.
- The maximum high-priority packet rate that is protected under congestion varies based on the platform and the maximum low-priority packet rate.

**CAUTION!** In a tunneled packet scenario with inner header parsing capability supported, such as IPv4-in-IPv6, IPv6-in-IPv4, or GRE, Ingress QoS for Layer 3 and Layer 4 fields works on inner header information rather than outer header information.

### 3.1 CPE Behavior — No Congestion State

Under normal conditions, the CPE behaves the same as when Ingress QoS is disabled except that a few cycles are spent finding the current congestion state and packets' Ingress QoS priority.

### 3.2 CPE Behavior — Congestion State

When the CPE is under congestion state, it either tries to free up processing cycles by dropping the low priority packets early during the RX processing stage, or the low priority packet processing is delayed to provide processing cycles to high priority traffic, in other words, by putting the Hi/Lo priority packets in different CPU queues.

Initially, when the CPE just enters the CPU congestion state, it will forward both high- and low-priority packets with some of the low priority packets dropped by Ingress QoS. Once the congestion state is reached as the congestion increases the low priority packet drop rate, but the high priority traffic is still protected. If the input packet rate keeps increasing there comes a point where all the low priority packets are dropped and if you go a little further, even the high priority traffic starts getting dropped.

### 3.3 Ingress QoS Packet Priority

Each received packet is assigned a low or high priority. A packet is assigned a high priority if meets one or more of the following criteria, otherwise it is assigned a low priority. This configuration is also relayed to the hardware network processor. If supported, the network processor will place packet in the proper receive queue based on the configured parameter when an Ingress QoS key hit occurs. (Criteria may vary per platform.)

- All multicast packets.
- Certain access level control packets, such as ARP, PPPOE Discovery and Session, and 802.1AG.
- SIP, RTSP, MGCP control packets (configured by VOIP).
- All of the RTP/RTCP connections established by SIP, RTSP, MGCP ALGs are also assigned high priority by adding the destination L4 (TCP/UDP) ports to the Ingress QoS table.
- Any entry added as high priority using Ingress QoS APIs.
- The ports that are added by default at initialization time or when an ALG is loaded are shown in [Table 2](#).

**Table 2: Default Entries Based on TCP/UDP + L4 Destination Ports**

UDP/TCP	Destination Port	Comment
TCP	80, 8080	HTTP
	53	DNS
UDP	53	DNS
	67, 68, 546, 547	DHCP
	2427, 2727	MGCP

**NOTE:** By default all the L4 packets are assigned a low priority unless it has been assigned a high priority by adding the L4 destination port to the Ingress QoS table.

**CAUTION!** In routing mode firewall should be enabled to load the ALGs (SIP, RTSP, etc.).

**CAUTION!** In bridge mode the ALGs (SIP, RTSP, etc.) are not loaded therefore the voice and video connections will be treated as low priority by default. To treat them as high priority, the voice and video applications should add the L4 destination ports to the ingress QoS port table using the Ingress QoS CLI or APIs.

### 3.4 Force Trap

If a network processor is capable of supporting trapping certain types of packets early, then the configuration will be relayed to the network processor. When a packet is classified as *force trap*, it is classified as high priority into the CPU received queue.

**NOTE:** This feature is only supported on Runner-based platforms. The types of traffic that are supported are:

- ARP
- PPPOE Discovery
- PPPOE session
- PTP 1588
- DHCP
- 802.1ag



## 3.5 Drop

Drop can be forced with a key match. The drop can also be relayed to the network processor. In the case of Runner-based devices, if the drop can be done in the classification stage, it will be done early, or else the drop happens before putting the packet into the receive queue.

Broadcom Confidential

## Chapter 4: CLI Commands

To see the list of Ingress QoS CLI commands just type **iq** at the shell prompt. (Commands vary per platform. The following display is a combination of all platforms. If a certain command is platform-specific, it will be mentioned in the later detailed section). Endianess for all the values should be in host order.

**NOTE:** “iq” is an alias for the “iqctl” CLI command.

```
# iq
Ingress QoS Control Utility:

::: Usage:

::: Ingress QoS SW System :
  iq status [-u]
    -u: new userspace status dump.
  iq enable
  iq disable
  iq flush
  iq addport
    --proto <0|1> --dport <1..65534> --ent <0|1> --prio <0|1>
    proto: 0 = TCP, 1 = UDP
    ent: 0 = dynamic, 1 = static
    prio: 0 = low, 1 = high
  iq remport --proto <0|1> --dport <1..65534> --ent <0|1>
    proto: 0 = TCP, 1 = UDP
    ent: 0 = dynamic, 1 = static
  iq getport --proto <0|1> --dport <1..65534>
    proto: 0 = TCP, 1 = UDP
  iq porttbl [-k]
    -k: dump from kernel
  iq addkeymask [field] --maskprio <0..15>
  iq remkeymask [field]
  iq keymasktbl
  iq addkey [field | action | attribute]
  iq remkey [field | action]
  iq getkey [field]
  iq keytbl
    field:
      --srcmac: MAC address in x:x:x:x:x:x format
      --dstmac: MAC address in x:x:x:x:x:x format
      --ethtype: ether type
      --outervid: Outer VLAN ID
      --outerpbit: Outer VLAN PBit
      --innervid: Inner VLAN ID
      --innerpbit: Inner VLAN PBit
      --l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
      --ipproto: such as TCP (6) or UDP (17)
      --srcip: Source IPv4/6 Address]
      --dstip: Destination IPv4/6 Address
      --dscp: 16-bit IP->DSCP value
      --ipv6flowlabel: 16-bit IPv6->flowlabel value
      --srcport: 16-bit L4 SRC Port
      --dstport: 16-bit L4 DST Port
    action:
      --prio: 0 = low, 1 = high
      --drop: drop packet
      --trap: trap packet before flow lookup
    attribute:
      --ent: 0 = dynamic, 1 = static
      --maskprio: <0..15> priority value, larger value = higher priority
  iq setdefaultprio --prototype <0> --protoval <0..255> --prio <0|1>
    prototype: 0 = ipproto
    protoval: protocol value (0 to 255)
    prio: 0 = low, 1 = high
  iq remdefaultprio --prototype <0> --protoval <0..255>
    prototype: 0 = ipproto
    protoval: protocol value (0 to 255)
```

## 4.1 Status

**Description** This command displays the current status of Ingress QoS feature and related information. (Default is printing from kernel. [-u] option will make it print from userspace, but with limited information.) The fields described in [Table 3](#) below are per interface for FAP-based platform.

**Syntax** `#iq status`  
**(FAP)** `[NTC iq] iq_get_status: Ingress QoS status : enabled`

```

-----IQ Status-----
      dev chnl loThr hiThr  used    dropped    cong
-----
      ENET   0   396   450     0         0         0
      ENET   1   528   600     0         0         0
      XTM    0    84    96     0         0         0
      XTM    1    33    37     0         0         0
FAP ENET    0   396   450     0         0         0
FAP XTM     0   132   150     0         0         0
FAP XTM     1    10    12     0         0         0

```

**Table 3: Status Fields for FAP-based (Per Interface)**

Field	Description
FAP	Indicates if the device or channel is managed by FAP. If this field is blank it means the device or channel is managed by the host.
dev	RX interface or device.
chnl	A channel on the RX interface also referred as queue interchangeably.
loThr	RX queue low threshold. When the queue depth becomes less than loThr, CPU congestion is removed for this queue.
hiThr	RX queue high threshold. When the queue depth becomes more than hiThr, CPU congestion is declared for this queue.
used	This field shows how many RXBDs for the channel are used at that instant. After the traffic stops, this field should be 0.
dropped	Number of packets dropped by Ingress QoS for the RX queue because of the CPU congestion. The reasons for higher packet drop may be either the RX queue (ring) size is small, or the input packet rate is higher than CPE can handle.
cong	This field shows if any of the RX queues is experiencing congestion. This field displays separate congestion status for host and FAP queues in hex format. The same value is displayed for all the host entries, and similarly same FAP congestion status for all the FAP entries. The <code>cong</code> field is interpreted as given below. Bits [31:10] = For future use. Bits [9:8] = One bit for each CMF FWD RX channel. Bit-8 for channel-0, bit-9 for channel-1, ... Bits [7:4] = One bit for each XTM RX channel. Bit-4 for channel-0, bit-5 for channel-1, ... Bits [3:0] = One bit for each Ethernet RX channel. Bit-0 for channel-0, bit-1 for channel-1, ...

**Syntax** Runner-based has up to 16 CPU RX queues, the output is based on QID.  
**(Runner)** # iq status  
 [NTC iq] iq\_get\_status: Ingress QoS status : enabled

```

-----IQ Status-----
qidx  rxpkts  queued  dropped  interrupts
-----
0      0      0      0      0
1      0      0      0      0
2      0      0      0      0
3      55     0      0      27
4      6678   0      0      3481
5      0      0      0      0
6      0      0      0      0
7      0      0      0      0
8      0      0      0      0
9      0      0      0      0
10     0      0      0      0
11     0      0      0      0
12     0      0      0      0
13     0      0      0      0

```

**Syntax** # iq status -u  
 BCM Ingress QoS Status: enabled

## 4.2 Enable

**Description** This command enables the Ingress QoS feature.

**Syntax** # iq enable

## 4.3 Disable

**Description** This command disables the Ingress QoS feature.

**Syntax** # iq disable

## 4.4 Flush

**Description** This command flushes all the dynamic entries.

**Syntax** # iq flush

## 4.5 Add Port

**Description** This command adds L4 destination port to the Ingress QoS port table. The fields are defined in [Table 4](#)

**Syntax** `# iq addport  
--proto <0|1> --dport <1..65534> --ent <0|1> --prio <0|1>`

**Table 4: Add Port Fields**

Field	Description
proto	UDP or TCP: 0 = TCP 1 = UDP
dport	L4 destination port, whose packet priority needs to be added.
ent	Entry type 0: Dynamic entry 1: Static entry
prio	Ingress QoS priority assigned to all the packets received with proto:dport combination 0 = Low 1 = High

## 4.6 Remove Port

**Description** This command removes L4 destination port from the Ingress QoS port table.

**Syntax** `# iq remport  
--proto <0|1> --dport <1..65534> --ent <0|1>`

## 4.7 Get Port

**Description** This command displays L4 destination port's entry type and packet priority from the Ingress. QoS port table. If a L4 destination port is not present in Ingress QoS port table, an entry type=0 (dynamic) and prio=0 (low) is returned.

**Syntax** `# iq getport  
--proto <0|1> --dport <1..65534>`

## 4.8 Dump Key Mask and Key Table

**Description** This command dumps the key mask and key entries of the Ingress QoS table for the requested protocol. Option -k can be used to dump the same tables in a different format from kernel.

**Syntax**

```
# iq porttbl

BCM Ingress QoS
  Status: enabled
  Total Static Entry Count: 16
  Total Dynamic Entry Count: 0

keymask table:
  index#0: prio = 14, refcnt = 4
           mask = --ethtype
  index#1: prio = 10, refcnt = 2
           mask = --ipproto --srcport
  index#2: prio = 9, refcnt = 0
           mask = --ipproto
  index#3: prio = 8, refcnt = 10
           mask = --ipproto --dstport

key table:
  index#60: static, refcnt = 1, SW hitcnt = 0
            field+value = --ethtype 0x8864
            action = prio = 1
  index#101: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 68
             action = prio = 1
  index#154: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 547
             action = prio = 1
  index#171: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 6 --dstport 8080
             action = prio = 1
  index#206: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 2727
             action = prio = 1
  index#231: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 2427
             action = prio = 1
  index#252: static, refcnt = 1, SW hitcnt = 0
             field+value = --ethtype 0x0806
             action = prio = 1
  index#269: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 53
             action = prio = 1
  index#289: static, refcnt = 1, SW hitcnt = 74078
             field+value = --ipproto 6 --dstport 80
             action = prio = 1
  index#380: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --srcport 53
             action = prio = 1
  index#399: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 6 --srcport 53
             action = prio = 1
  index#401: static, refcnt = 1, SW hitcnt = 0
             field+value = --ethtype 0x8902
             action = prio = 1
  index#427: static, refcnt = 1, SW hitcnt = 0
             field+value = --ipproto 17 --dstport 546
             action = prio = 1
  index#428: static, refcnt = 1, SW hitcnt = 0
             field+value = --ethtype 0x8863
             action = prio = 1
  index#498: static, refcnt = 1, SW hitcnt = 1
             field+value = --ipproto 17 --dstport 67
             action = prio = 1
  index#510: static, refcnt = 1, SW hitcnt = 1
             field+value = --ipproto 6 --dstport 53
             action = prio = 1
```

**Syntax**

# iq porttbl -k

key mask table

index	priority	key mask	key option	ref_cnt
00000000	00000014	0x00000008	0x00000001	4
00000001	00000010	0x00008400	0x00000001	2
00000002	00000009	0x00000400	0x00000001	0
00000003	00000008	0x00010400	0x00000001	10

key hash table

Total static Entry Count: 16

Total dynamic Entry Count: 0

index	key mask	key option	action	ref_cnt	hits	static
field = value						
00000060	0x00000008	0x00000001	0x00000047	0000001	000000	yes
EtherType = 0x8864						
00000101	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 68						
00000154	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 547						
00000171	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x06						
DST Port = 8080						
00000206	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 2727						
00000231	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 2427						
00000252	0x00000008	0x00000001	0x00000047	0000001	000000	yes
EtherType = 0x0806						
00000269	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 53						
00000289	0x00010400	0x00000001	0x00000047	0000001	074330	yes
IP PROTO = 0x06						
DST Port = 80						
00000380	0x00008400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
SRC Port = 53						
00000399	0x00008400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x06						
SRC Port = 53						
00000401	0x00000008	0x00000001	0x00000047	0000001	000000	yes
EtherType = 0x8902						
00000427	0x00010400	0x00000001	0x00000047	0000001	000000	yes
IP PROTO = 0x11						
DST Port = 546						
00000428	0x00000008	0x00000001	0x00000047	0000001	000000	yes
EtherType = 0x8863						
00000498	0x00010400	0x00000001	0x00000047	0000001	000001	yes
IP PROTO = 0x11						
DST Port = 67						
00000510	0x00010400	0x00000001	0x00000047	0000001	000001	yes
IP PROTO = 0x06						
DST Port = 53						

## 4.9 Add Key Mask

**Description** This command adds key mask. A key mask must be added first before a key using the field set in the key mask (fields can be ORed) is added. **Ingress QoS must be flushed with no dynamic entries left and disabled before adding a key mask.** A unique priority value (from 0 to 15) must be configured with `--maskprio`. Adding a key mask will fail if attempts to add with a priority that exists in Ingress QoS.

**Syntax** `# iq addkeymask [fields w/o value] -maskprio <0..15>`

(Runner)

```
field:
  --srcmac: MAC address in x:x:x:x:x:x format
  --dstmac: MAC address in x:x:x:x:x:x format
  --ethtype: ether type
  --outervid: Outer VLAN ID
  --outerpbit: Outer VLAN PBit
  --innervid: Inner VLAN ID
  --innerpbit: Inner VLAN PBit
  --l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
  --ipproto: such as TCP (6) or UDP (17)
  --srcip: Source IPv4/6 Address
  --dstip: Destination IPv4/6 Address
  --dscp: 6-bit IP->DSCP value
  --ipv6flowlabel: 20-bit IPv6->flowlabel value
  --srcport: 16-bit L4 SRC Port
  --dstport: 16-bit L4 DST Port
attribute:
  --maskprio: <0..15> priority value, larger value = higher priority
```

## 4.10 Remove Key Mask

**Description** This command removes key mask.

**Syntax** `# iq remkeymask [fields w/o value]`

(Runner)

```
field:
  --srcmac: MAC address in x:x:x:x:x:x format
  --dstmac: MAC address in x:x:x:x:x:x format
  --ethtype: ether type
  --outervid: Outer VLAN ID
  --outerpbit: Outer VLAN PBit
  --innervid: Inner VLAN ID
  --innerpbit: Inner VLAN PBit
  --l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
  --ipproto: such as TCP (6) or UDP (17)
  --srcip: Source IPv4/6 Address
  --dstip: Destination IPv4/6 Address
  --dscp: 6-bit IP->DSCP value
  --ipv6flowlabel: 20-bit IPv6->flowlabel value
  --srcport: 16-bit L4 SRC Port
  --dstport: 16-bit L4 DST Port
```



## 4.11 Dump Key Mask Table

**Description** This command displays key mask table.

**Syntax**  
(Runner)

```
# iq keymasktbl
keymask table:
    index#0: prio = 14, refcnt = 4
              mask = --ethtype
    index#1: prio = 10, refcnt = 2
              mask = --ipproto --srcport
    index#2: prio = 9, refcnt = 0
              mask = --ipproto
    index#3: prio = 8, refcnt = 10
              mask = --ipproto --dstport
```

## 4.12 Add Key

**Description** This command adds key. Default action is nop. Entry will be created as dynamic by default.

**Syntax**  
(Runner)

```
# iq addkey [field | action | attribute]
```

**field:**

- srcmac: MAC address in x:x:x:x:x:x format
- dstmac: MAC address in x:x:x:x:x:x format
- ethtype: ether type
- outervid: Outer VLAN ID
- outerpbit: Outer VLAN PBit
- innervid: Inner VLAN ID
- innerpbit: Inner VLAN PBit
- l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
- ipproto: such as TCP (6) or UDP (17)
- srcip: Source IPv4/6 Address
- dstip: Destination IPv4/6 Address
- dscp: 6-bit IP->DSCP value
- ipv6flowlabel: 20-bit IPv6->flowlabel value
- srcport: 16-bit L4 SRC Port
- dstport: 16-bit L4 DST Port

**action:**

- prio: 0 = low, 1 = high
- drop: drop packet
- trap: trap packet before flow lookup

**attribute:**

- ent: 0 = dynamic, 1 = static

## 4.13 Remove Key

**Description** This command removes key.

**Syntax** # iq remkey [field | action]

(Runner)

```
field:
--srcmac: MAC address in x:x:x:x:x:x format
--dstmac: MAC address in x:x:x:x:x:x format
--ethtype: ether type
--outervid: Outer VLAN ID
--outerpbit: Outer VLAN PBit
--innervid: Inner VLAN ID
--innerpbit: Inner VLAN PBit
--l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
--ipproto: such as TCP (6) or UDP (17)
--srcip: Source IPv4/6 Address
--dstip: Destination IPv4/6 Address
--dscp: 6-bit IP->DSCP value
--ipv6flowlabel: 20-bit IPv6->flowlabel value
--srcport: 16-bit L4 SRC Port
--dstport: 16-bit L4 DST Port
action:
--prio: 0 = low, 1 = high
--drop: drop packet
--trap: trap packet before flow lookup
```

## 4.14 Get Key

**Description** This command gets key based on the field and prints out info about this key.

**Syntax** # iq getkey [field]

(Runner)

```
field:
--srcmac: MAC address in x:x:x:x:x:x format
--dstmac: MAC address in x:x:x:x:x:x format
--ethtype: ether type
--outervid: Outer VLAN ID
--outerpbit: Outer VLAN PBit
--innervid: Inner VLAN ID
--innerpbit: Inner VLAN PBit
--l3proto: e.g., IPv4 (0x0800) or IPv6 (0x86DD)
--ipproto: such as TCP (6) or UDP (17)
--srcip: Source IPv4/6 Address
--dstip: Destination IPv4/6 Address
--dscp: 6-bit IP->DSCP value
--ipv6flowlabel: 20-bit IPv6->flowlabel value
--srcport: 16-bit L4 SRC Port
--dstport: 16-bit L4 DST Port
```

## 4.15 Dump Key Table

**Description** This command prints all the keys.

**Syntax** # iq keytbl  
**(Runner)** key table:

```
index#60: static, refcnt = 1, SW hitcnt = 0
        field+value = --ethtype 0x8864
        action = prio = 1
index#101: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 68
        action = prio = 1
index#154: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 547
        action = prio = 1
index#171: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 6 --dstport 8080
        action = prio = 1
index#206: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 2727
        action = prio = 1
index#231: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 2427
        action = prio = 1
index#252: static, refcnt = 1, SW hitcnt = 0
        field+value = --ethtype 0x0806
        action = prio = 1
index#269: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 53
        action = prio = 1
index#289: static, refcnt = 1, SW hitcnt = 75218
        field+value = --ipproto 6 --dstport 80
        action = prio = 1
index#380: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --srcport 53
        action = prio = 1
index#399: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 6 --srcport 53
        action = prio = 1
index#401: static, refcnt = 1, SW hitcnt = 0
        field+value = --ethtype 0x8902
        action = prio = 1
index#427: static, refcnt = 1, SW hitcnt = 0
        field+value = --ipproto 17 --dstport 546
        action = prio = 1
index#428: static, refcnt = 1, SW hitcnt = 0
        field+value = --ethtype 0x8863
        action = prio = 1
index#498: static, refcnt = 1, SW hitcnt = 1
        field+value = --ipproto 17 --dstport 67
        action = prio = 1
index#510: static, refcnt = 1, SW hitcnt = 4
        field+value = --ipproto 6 --dstport 53
        action = prio = 1
```

## 4.16 Set Default Priority for IP Protocol Value

**Description** This command configures the default priority value given an IP Protocol value.

**Syntax** # iq setdefaultprio --prototype <0> --protoval <0..255> --prio <0/1>  
(FAP) prototype: 0 = IPPROTO  
protoval: protocol value (0 to 255)  
prio: 0 = low, 1 = high

## 4.17 Remove Configured Default Priority of an IP Protocol Value

**Description** This command removes the previously configured default priority set on an IP Protocol Value.

**Syntax** # iq remdefaultprio --prototype <0> --protoval <0..255>  
(Runner) prototype: 0 = IPPROTO  
protoval: protocol value (0 to 255)

## 4.18 Example

The following example will add the required keymask and key entry to drop packet with outervid = 0x20 and srcip = 12.34.56.78.

```
# iq keymasktbl
keymask table:
  index#0: prio = 14, refcnt = 4
           mask = --ethtype
  index#1: prio = 10, refcnt = 2
           mask = --ipproto --srcport
  index#2: prio = 9, refcnt = 0
           mask = --ipproto
  index#3: prio = 8, refcnt = 10
           mask = --ipproto --dstport

# iq disable
# iq flush
# iq addkeymask --outervid --srcip --maskprio 13
# iq enable
# iq addkey --outervid 0x20 --srcip 12.34.56.78 --drop
# iq keymasktbl
keymask table:
  index#0: prio = 14, refcnt = 4
           mask = --ethtype
  index#1: prio = 13, refcnt = 1
           mask = --outervid --srcip
  index#2: prio = 10, refcnt = 2
           mask = --ipproto --srcport
  index#3: prio = 8, refcnt = 10
           mask = --ipproto --dstport

# iq getkey --outervid 0x20 --srcip 12.34.56.78
  index#0: dynamic, refcnt = 0, SW hitcnt = 0
           field+value = --outervid 32 --srcip 12.34.56.78
           action = drop
```

## 4.19 Example

The following example will modify the default ARP packet trapped to high priority CPU receive queue action to a new early CPU trap action for all ARP packets.

```
# iq getkey --ethtype 0x0806
    index#0: static, refcnt = 0, SW hitcnt = 0
            field+value = --ethtype 0x0806
            action = prio = 1
# iq remkey --ethtype 0x0806
# iq addkey --ethtype 0x0806 --trap --ent 1
# iq getkey --ethtype 0x0806
    index#0: static, refcnt = 0, SW hitcnt = 0
            field+value = --ethtype 0x0806
            action = trap
```

Broadcom Confidential

## Chapter 5: APIs

Ingress QoS exports APIs for the user in kernel space, such as drivers or modules, to interact and configure Ingress QoS. Ingress QoS API prototypes are defined in CommEngine/kernel/linux/include/linux/iqos.h. Endianness for all the values should be in host order.

### 5.1 Add L4 Destination Port

<b>Description</b>	This API adds a L4 destination port, protocol, packet priority, and the entry type to the Ingress QoS table. (Legacy API, it is backward compatible.)								
<b>Syntax</b>	<pre>int iqos_add_L4port     (uint8_t ipProto, uint16_t destPort, iqos_ent_t ent, iqos_prio_t prio);</pre>								
<b>Parameters</b>	<table> <tr> <td>ipProto</td><td>IP L4 protocol TCP or UDP</td></tr> <tr> <td>destPort</td><td>L4 destination port.</td></tr> <tr> <td>ent</td><td>Entry type: static or dynamic.</td></tr> <tr> <td>prio</td><td>Packet priority assigned to the incoming packet.</td></tr> </table>	ipProto	IP L4 protocol TCP or UDP	destPort	L4 destination port.	ent	Entry type: static or dynamic.	prio	Packet priority assigned to the incoming packet.
ipProto	IP L4 protocol TCP or UDP								
destPort	L4 destination port.								
ent	Entry type: static or dynamic.								
prio	Packet priority assigned to the incoming packet.								
<b>Return Value</b>	Success: 0 Failure: Otherwise								

### 5.2 Remove L4 Destination Port

<b>Description</b>	This API removes a previously added L4 destination port, protocol, and the entry type from the Ingress QoS table. (Legacy API, it is backward compatible.)						
<b>Syntax</b>	<pre>int iqos_rem_L4port(uint8_t ipProto, uint16_t destPort, iqos_ent_t ent);</pre>						
<b>Parameters</b>	<table> <tr> <td>ipProto</td><td>IP L4 protocol TCP or UDP</td></tr> <tr> <td>destPort</td><td>L4 destination port.</td></tr> <tr> <td>ent</td><td>Entry type: static or dynamic</td></tr> </table>	ipProto	IP L4 protocol TCP or UDP	destPort	L4 destination port.	ent	Entry type: static or dynamic
ipProto	IP L4 protocol TCP or UDP						
destPort	L4 destination port.						
ent	Entry type: static or dynamic						
<b>Return Value</b>	Success: 0 Failure: Otherwise						

### 5.3 Get L4 Destination Port Priority

<b>Description</b>	This API returns the current assigned packet priority for a L4 destination port, and protocol from the Ingress QoS table. (Legacy API, it is backward compatible.)				
<b>Syntax</b>	<pre>int iqos_prio_L4port(uint8_t ipProto, uint16_t destPort);</pre>				
<b>Parameters</b>	<table> <tr> <td>ipProto</td><td>IP L4 protocol TCP or UDP</td></tr> <tr> <td>destPort</td><td>L4 destination port</td></tr> </table>	ipProto	IP L4 protocol TCP or UDP	destPort	L4 destination port
ipProto	IP L4 protocol TCP or UDP				
destPort	L4 destination port				
<b>Return Value</b>	Success: packet priority IQOS_PRIO_LOW/IQOS_PRIO_HIGH Failure: packet priority IQOS_PRIO_LOW				

## 5.4 Initialize Keymask Parameter

**Description** This API is used to initialize parameter holder for configuring Ingress QoS Keymask.

**Syntax** `int iqos_keymask_param_start(iqos_param_t *param);`

**Parameters** param Ingress QoS parameter holder

**Return Value** Success: 0, parameter holder is initialized and set to be used for Keymask operation.  
Failure: Otherwise

## 5.5 Set Field and Value to the Keymask Parameter

**Description** This API sets field and its value to the Ingress QoS parameter holder for configuring Ingress QoS Keymask.

**Syntax** `int iqos_keymask_param_field_set(iqos_param_t *param, uint32_t field, uint32_t *val_ptr);`

**Parameters** param Ingress QoS parameter holder  
field Ingress QoS field type (details of field type and value are in later table)  
val\_ptr value holder

**Return Value** Success: 0  
Failure: Otherwise

## 5.6 Commit and Add the Keymask Parameter

**Description** This API commits the Ingress QoS parameter holder for a Keymask add event.

**Syntax** `int iqos_keymask_commit_and_add(iqos_param_t *param, uint8_t prio);`

**Parameters** param Ingress QoS parameter holder  
prio Priority

**Return Value** Success: 0  
Failure: Otherwise

## 5.7 Commit and Delete the Keymask Parameter

**Description** This API commits the Ingress QoS parameter holder for a Keymask delete event.

**Syntax** `int iqos_keymask_commit_and_delete(iqos_param_t *param);`

**Parameters** param Ingress QoS parameter holder

**Return Value** Success: 0  
Failure: Otherwise

## 5.8 Initialize Key Parameter

**Description** This API is used to initialize parameter holder for configuring Ingress QoS Key.

**Syntax** `iqos_key_param_start(iqos_param_t *param);`

**Parameters** param Ingress QoS parameter holder

**Return Value** Success: 0, parameter holder is initialized and set to be used for Key operation.  
Failure: Otherwise

## 5.9 Set Field and Value to the Key Parameter

**Description** This API sets field and its value to the Ingress QoS parameter holder for configuring Ingress QoS Key.

**Syntax** `int iqos_key_param_field_set(iqos_param_t *param, uint32_t field, uint32_t *val_ptr, uint32_t val_size);`

**Parameters** param Ingress QoS parameter holder  
field Ingress QoS field type (details of field type and value are in later table)  
val\_ptr Value holder  
val\_size Size of data in val\_ptr

**Return Value** Success: 0  
Failure: Otherwise

## 5.10 Set Action and Value to the Key Parameter

**Description** This API sets action and its value to the Ingress QoS parameter holder for configuring Ingress QoS Key.

**Syntax** `int iqos_key_param_action_set(iqos_param_t *param, uint32_t action, uint32_t value);`

**Parameters** param Ingress QoS parameter holder  
action Ingress QoS action type (details of action type and value are in later table)  
value Value used for the action

**Return Value** Success: 0  
Failure: Otherwise

## 5.11 Commit and Add the Key Parameter

**Description** This API commits the Ingress QoS parameter holder for a Key add event.

**Syntax** `int iqos_key_commit_and_add(iqos_param_t *param, uint8_t type);`

**Parameters** param Ingress QoS parameter holder  
type dynamic (0) or static (1)

**Return Value** Success: 0  
Failure: Otherwise



## 5.12 Commit and Delete the Key Parameter

**Description** This API commits the Ingress QoS parameter holder for a Key delete event.

**Syntax** `int iqos_key_commit_and_delete(iqos_param_t *param, uint8_t type);`

**Parameters**

param	Ingress QoS parameter holder
type	Dynamic (0) or static (1)

**Return Value** Success: 0  
Failure: Otherwise

## 5.13 Commit and Get the Key Parameter

**Description** This API commits the Ingress QoS parameter holder for a Key get event.

**Syntax** `int iqos_key_commit_and_get(iqos_param_t *param);`

**Parameters**

param	Ingress QoS parameter holder
-------	------------------------------

**Return Value** Success: 0 when entry is found, and update action and its value in the parameter holder  
Failure: Otherwise

## 5.14 Flush All Dynamic Key Entries

**Description** This API flushes all dynamic key entries and key mask if there is no key referred to the specific key mask.

**Syntax** `int iqos_flush(void);`

**Parameters** None

**Return Value** None

## 5.15 Set Status

**Description** This API flushes all dynamic key entries and key mask if there is no key referred to the specific key mask.

**Syntax** `int iqos_set_status(uint32_t status);`

**Parameters** Status      0 = disable, 1 = enable

**Return Value** Success: 0  
Failure: Otherwise

**Table 5: Fields Table**

Field	Description
IQOS_FIELD_SRC_MAC	Source MAC Address. If used for configuring key, val_ptr must point to starting of a 6 bytes memory and val_size given should be 6.
IQOS_FIELD_DST_MAC	Destination MAC Address. If used for configuring key, val_ptr must point to starting of a 6 bytes memory and val_size given should be 6.
IQOS_FIELD_ETHER_TYPE	Ether_type. 2-byte EtherType value in the Ethernet header.
IQOS_FIELD_OUTER_VID	Outer VLAN ID. 2-byte VLAN ID value in the outer VLAN tag. (Also used in single VLAN tag scenario.)
IQOS_FIELD_OUTER_PBIT	Outer PBIT. 1-byte PBIT value in the outer VLAN tag. (Also used in single VLAN tag scenario.)
IQOS_FIELD_INNER_VID	Inner VLAN ID. 2-byte VLAN ID value in the inner VLAN tag. (Used in double-tagged scenario. There is no support for 3 or more VLAN tags.)
IQOS_FIELD_INNER_PBIT	Inner PBIT. 1-byte PBIT value in the inner VLAN tag. (Used in double-tagged scenario. There is no support for 3 or more VLAN tags.)
IQOS_FIELD_L3_PROTO	L3 Protocol. 2-byte L3 protocol value, i.e., IPv4 = 0x8000.
IQOS_FIELD_IP_PROTO	IP Protocol. 1-byte IP protocol value, i.e., TCP = 6.
IQOS_FIELD_SRC_IP	Source IP address. 4-byte IP address for IPv4 or 16-byte for IPv6.
IQOS_FIELD_DST_IP	Destination IP address. 4-byte IP address for IPv4 or 16-byte for IPv6.
IQOS_FIELD_DSCP	DSCP. 1-byte DSCP value from IPv4 header.
IQOS_FIELD_IPV6_FLOW_LABEL	IPv6 Flow label. 4-byte Flow Label value from IPv6 header.
IQOS_FIELD_SRC_PORT	L4 Source Port. 2-byte source port from L4 (TCP/UDP) header.
IQOS_FIELD_DST_PORT	L4 Destination Port. 2-byte destination port from L4 (TCP/UDP) header.

**Table 6: Actions Table**

Action	Description
IQOS_ACTION_NOP	No operation.
IQOS_ACTION_PRIO	Setting the received queue priority when Ingress QoS key hits. Value is either 0 for low priority or 1 for high priority.
IQOS_ACTION_DROP	Dropping the packet when key hits.
IQOS_ACTION_TRAP	Trapping the packet early in classification to CPU when key hits. (If given key is supported by hardware network processor, it usually traps to high priority receive queue.)

## 5.16 Configuration Sequence

There are two important criteria that one has to follow:

1. A matching key mask must be configured before adding a key.
2. In order to add a new keymask, Ingress QoS must have no existing dynamic entry and it has to be in disabled stage.

The following shows an example of how to add a new key mask and a key to Ingress QoS from a kernel space user module.

```
Int rc;
iqos_param_t iqos_param;
uint8_t src_mac[6] = {0x12, 0x34, 0x56, 0x78, 0x9a, 0xbc};
uint32_t dst_ip = 0xc0a80120;
uint16_t dst_port = 123;

rc = iqos_set_status(0); /* to disable Ingress QoS */
if (rc)
    return rc;
iqos_flush();
rc = iqos_keymask_param_start(&iqos_param);
rc = rc? rc : iqos_keymask_param_field_set(&iqos_param, IQOS_FIELD_SRC_MAC, NULL);
rc = rc? rc : iqos_keymask_param_field_set(&iqos_param, IQOS_FIELD_DST_IP, NULL);
rc = rc? rc : iqos_keymask_param_field_set(&iqos_param, IQOS_FIELD_DST_PORT, NULL);
rc = rc? rc : iqos_keymask_commit_and_add(&iqos_param, 8);
if (rc)
    printk("error!\n");
rc = iqos_set_status(1); /* to enable Ingress QoS */

rc = rc? rc : iqos_key_param_start(&iqos_param);
rc = rc? rc : iqos_key_param_field_set(&iqos_param, IQOS_FIELD_SRC_MAC, (uint32_t *)src_mac, 6);
rc = rc? rc : iqos_key_param_field_set(&iqos_param, IQOS_FIELD_DST_IP, (uint32_t *)&dst_ip, 4);
rc = rc? rc : iqos_key_param_field_set(&iqos_param, IQOS_FIELD_DST_PORT, (uint32_t *)&dst_port, 2);
rc = rc? rc : iqos_key_commit_and_add(&iqos_param, 1); /* static entry */
```

## Chapter 6: Performance

Under normal conditions, with no CPU congestion, all packets are forwarded.

Under CPU congestion only high-priority packets are forwarded and all low priority packets are dropped by Ingress QoS.

### 6.1 All High-Priority Flows

If all of the flows are at the same priority, none of the flows gets priority. In this scenario, Ingress QoS tries to protect all high-priority flows. However, none of them are protected because there is no priority among them. Once congestion is reached, packets from all the flows are dropped proportionately because the accelerator/CPU cannot handle the packet rate.

Behavior in this scenario is very similar to when the Ingress QoS feature was not implemented (before 4.10 release) or the feature is disabled.

The overall performance should be slightly lower than when the feature is disabled.

### 6.2 All Low-Priority Flows

When all of the flows are of same priority (low) none of the flows get priority. In this scenario, Ingress QoS attempts to protect a high priority flow, but there are no high priority flows. After congestion, packets from all the low priority flows are dropped proportionately by the Ingress QoS to preserve CPU cycles for a high-priority flow.

**NOTE:** The maximum forwarding rate should be slightly lower than when the feature is disabled, but after congestion is hit the output rate should come down steeply with the increase in the low-priority input packet rate.

**CAUTION!** As the low-priority input rate increases, the low-priority output rate decreases after CPU congestion is reached. This is contrary to what most people expect and differs from the behavior when Ingress QoS feature was not implemented.

### 6.3 Mix of High- and Low-Priority Flows

In this scenario, Ingress QoS tries to protect a high priority flow at the cost of low-priority flows under CPU congestion. Packets from a low priority flow are dropped when there is CPU congestion. When there is no congestion all packets are forwarded. In addition, when a packet is classified as high priority by the network processor, it will be placed into high priority receive queue. The Ethernet driver will always serve the high priority queue first. This ensures the handling of high priority packets, but there could be cases of tail drop in the low priority receive queue.

### 6.4 Early Trap and Drop

(Applicable only for the Runner-based platform.) When using these two features, it will enable ingress filter in Runner. This ingress filter is executed on every packet in the critical data path in the Runner. Therefore, it has negative impact on performance of other traffic.

## Revision History

### 963XX-AN403-R, December 18, 2018

- Updated Introduction to [Chapter 1, Overview](#)
- Added Note to [Chapter 1, Overview](#)
- Updated Introduction to [Chapter 3, Ingress QoS](#)
- Added Caution note to Introduction of [Chapter 3, Ingress QoS](#)
- Updated content and changed heading to “CPE Behavior — No Congestion State” on page 7
- Updated content and changed heading to “CPE Behavior — Congestion State” on page 7
- Updated “[Force Trap](#)” on page 8
- Updated introduction to [Chapter 4, CLI Commands](#)
- Added Note to [Chapter 4, CLI Commands](#)
- Updated `l3proto` value for IPv6 throughout
- Updated “[Example](#)” on page 20
- Added “[Example](#)” on page 21
- Updated introduction to [Chapter 5, APIs](#)

### 963XX-AN402-R, July 11, 2018

- Updated introduction to Chapter 1, Overview
- Updated introduction to Chapter 2, Enable and Disable Ingress QoS
- Updated “Enable Ingress QoS” on page 6
- Updated “Disable Ingress QoS” on page 6
- Updated introduction to Chapter 3, Ingress QoS
- Updated “Ingress QoS Packet Priority” on page 8
- Updated Table 2, Default Entries Based on TCP/UDP + L4 Destination Ports
- Added “Force Trap” on page 8
- Added “Drop” on page 8
- Updated Chapter 4, CLI Commands
- Added “Status” on page 10
- Updated “Flush” on page 11
- Updated “Dump Key Mask and Key Table” on page 13
- Added “Add Key Mask” on page 15
- Added “Remove Key Mask” on page 15
- Added “Dump Key Mask Table” on page 16
- Added “Add Key” on page 16
- Added “Remove Key” on page 17
- Added “Get Key” on page 17
- Added “Dump Key Table” on page 18
- Added “Set Default Priority for IP Protocol Value” on page 19
- Added “Remove Configured Default Priority of an IP Protocol Value” on page 19
- Added “Example” on page 19
- Updated introduction to Chapter 5, APIs
- Updated “Add L4 Destination Port” on page 20
- Updated “Remove L4 Destination Port” on page 20

- Updated “Get L4 Destination Port Priority” on page 20
- Added “Set Field and Value to the Keymask Parameter” on page 21
- Added “Commit and Add the Keymask Parameter” on page 21
- Added “Commit and Delete the Keymask Parameter” on page 21
- Added “Set Field and Value to the Keymask Parameter” on page 21
- Added “Initialize Key Parameter” on page 22
- Added “Set Field and Value to the Key Parameter” on page 22
- Added “Set Action and Value to the Key Parameter” on page 22
- Added “Commit and Add the Key Parameter” on page 22
- Added “Commit and Delete the Key Parameter” on page 23
- Added “Commit and Get the Key Parameter” on page 23
- Added “Flush All Dynamic Key Entries” on page 23
- Added “Set Status” on page 24
- Updated “Mix of High- and Low-Priority Flows” on page 26
- Added “Early Trap and Drop” on page 26

## Previous Release History

Revision	Date	Change Description
963XX-AN101-R	01/18/15	<b>Updated:</b> <ul style="list-style-type: none"> <li>■ “Add Port” on page 12</li> <li>■ “Dump Port Table” on page 13</li> <li>■ “Add L4 Destination Port” on page 14</li> </ul>
963XX-AN400-R	12/14/10	Initial release

Broadcom, the pulse logo, Connecting everything, Avago Technologies, Avago, and the A logo are among the trademarks of Broadcom and/or its affiliates in the United States, certain other countries, and/or the EU.

Copyright © 2010–2018 Broadcom. All Rights Reserved.

The term “Broadcom” refers to Broadcom Inc. and/or its subsidiaries. For more information, please visit [www.broadcom.com](http://www.broadcom.com).

Broadcom reserves the right to make changes without further notice to any products or data herein to improve reliability, function, or design. Information furnished by Broadcom is believed to be accurate and reliable. However, Broadcom does not assume any liability arising out of the application or use of this information, nor the application or use of any product or circuit described herein, neither does it convey any license under its patent rights nor the rights of others.